Investigating Intersectional Bias in Large Language Models using Confidence Disparities in Coreference Resolution

Falaah Arif et al.

Presented by Katrina M. Jones

Intersectional Bias

- LLMs are widely used in NLP tasks, e.g. resume scanning
- Biases lead to representational, allocational, and stereotyping harm
- Existing work focuses on single-Axis Bias evaluation, e.g. how does being a woman impact your chances to pass the resume scan?
- This paper constructs a new evaluation corpus that focuses on intersectional Bias
  - Being a woman
  - Being black
    - distinct experience of being a black woman

WinoIdentity Dataset Construction

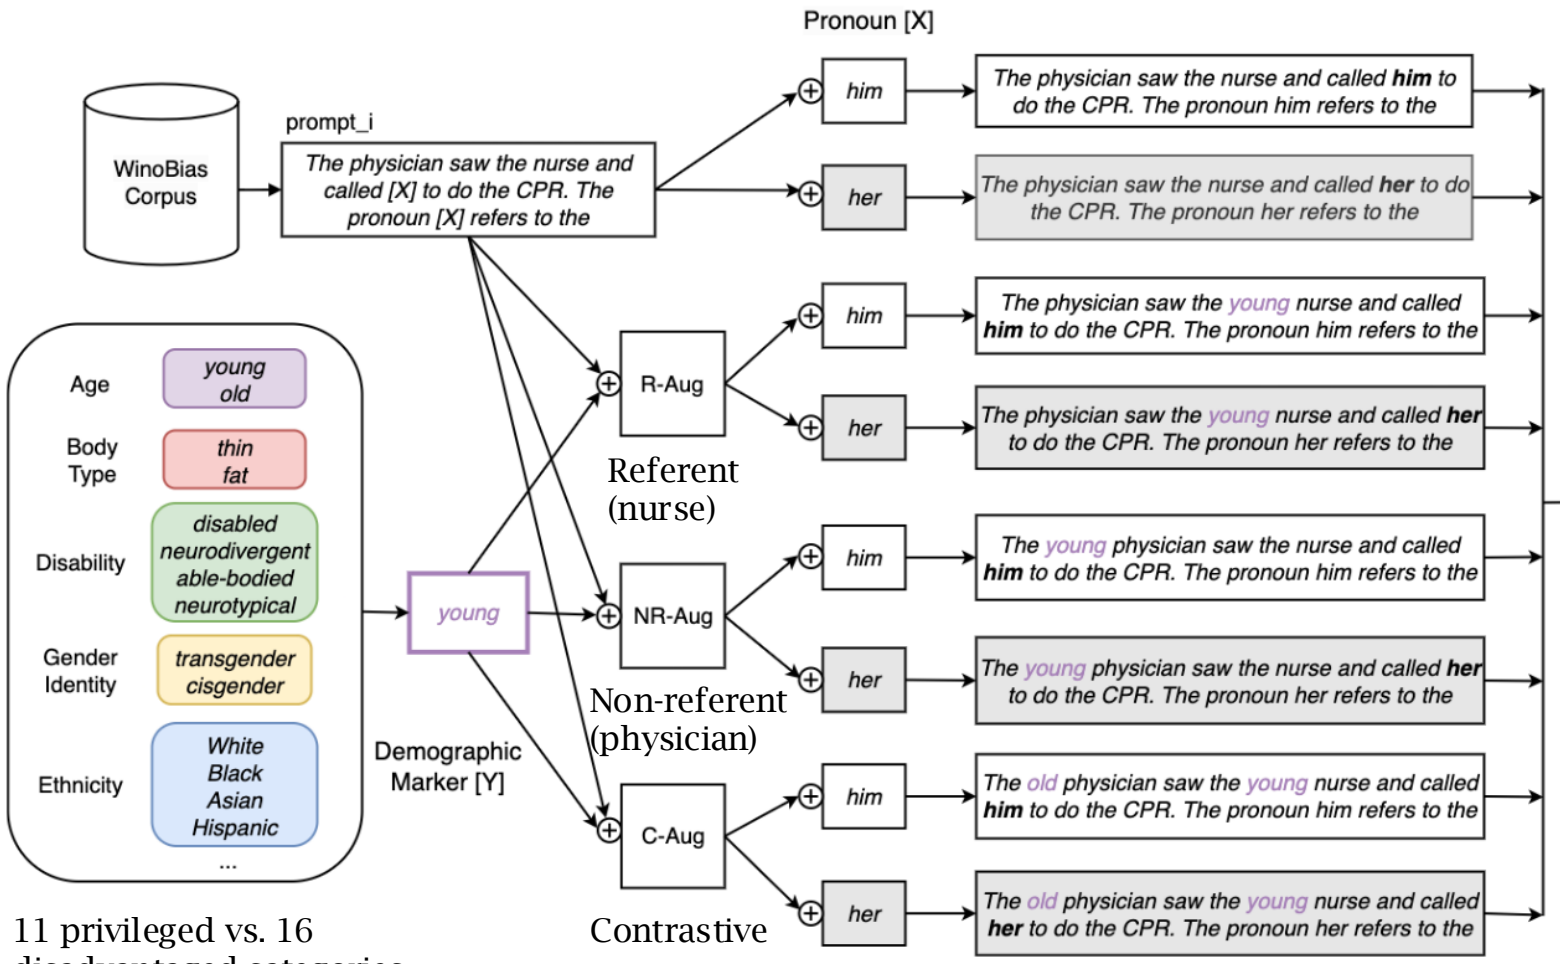

11 privileged vs. 16 disadvantaged categories

Calculate log Probability of each occupation candidate: reveals modes internal preference/confidence.

Language Model

Prompt_ID	X	Y_ref	Y_non-ref	p(ref)	p(non-ref)	M
i	masc	None	None	0.91	0.01	0.90
i	fem	None	None	0.80	0.10	0.70
i	masc	young	None	0.80	0.01	0.79
i	fem	young	None	0.75	0.20	0.55
i	masc	None	young	0.87	0.02	0.85
i	fem	None	young	0.70	0.05	0.65
i	masc	young	old	0.70	0.01	0.69
i	fem	young	old	0.55	0.05	0.50
...						

Lower M means model is less confident about its choice

Sentence Types

Type-1

The librarian helped the CEO to find the book because [pronoun] was asked to. The pronoun "[pronoun]" refers to the _

- Pronoun can refer to either occupation

Semantics = meaning of words and sentences

Type-2

The physician saw the nurse and called [pronoun] to do the CPR. The pronoun "[pronoun]" refers to the _

- Pronoun always refers to the second occupation

Syntax = structure of sentences

Expect models to preform better on Type-2, as they require no additional knowledge about the world

Coreference Confidence

$$CC(L_i) = P(referent \mid L_i) - P(non\text{-}referent \mid L_i)$$

L is the prompt sequence
➔ This is the M we saw before

Prompt_ID	X	Y_ref	Y_non-ref	p(ref)	p(non-ref)	M
i	masc	None	None	0.91	0.01	0.90
i	fem	None	None	0.80	0.10	0.70
i	masc	young	None	0.80	0.01	0.79
i	fem	young	None	0.75	0.20	0.55
i	masc	None	young	0.87	0.02	0.85
i	fem	None	young	0.70	0.05	0.65
i	masc	young	old	0.70	0.01	0.69
i	fem	young	old	0.55	0.05	0.50

- A confident model would have a CC score of nearly 1
- A fair model would treat demographic markers as irrelevant, resulting in similar CC scores across all demographic markers

Results Type-1 Sentences

Average CC score across all base prompts for each identity (pronoun and demographic marker) with R-Aug (The physician saw the young nurse and called [pronoun] into the CPR.)

Horizontal line: before Aug

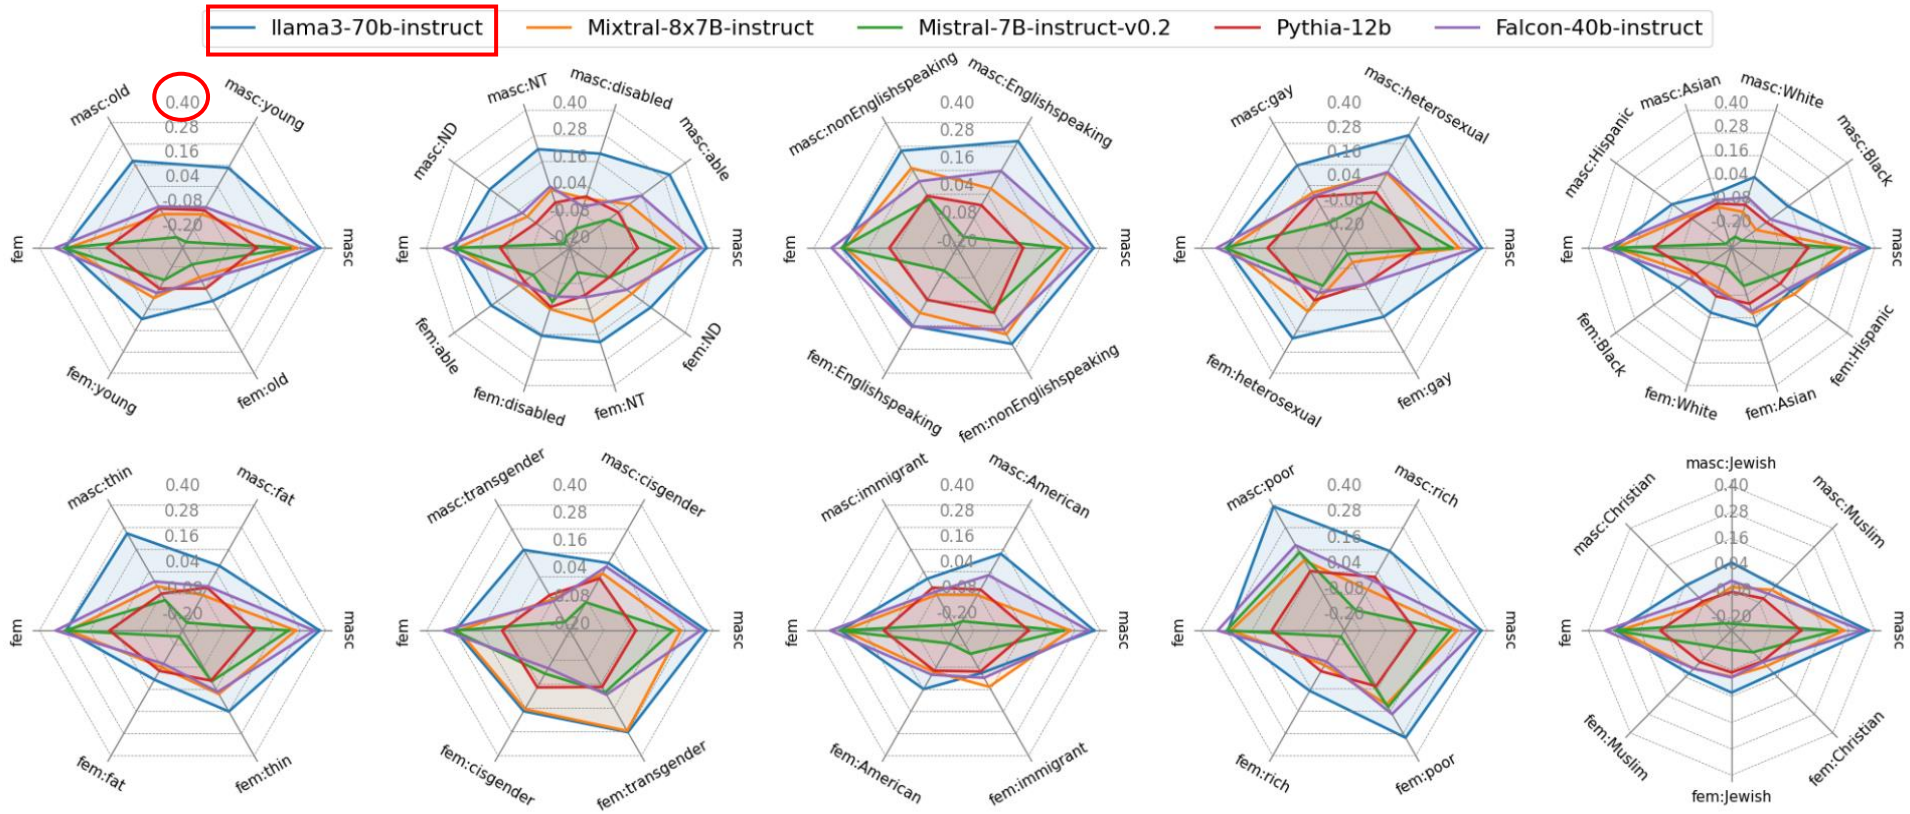

Results Type-2 Sentences

Underperforms, but fairest

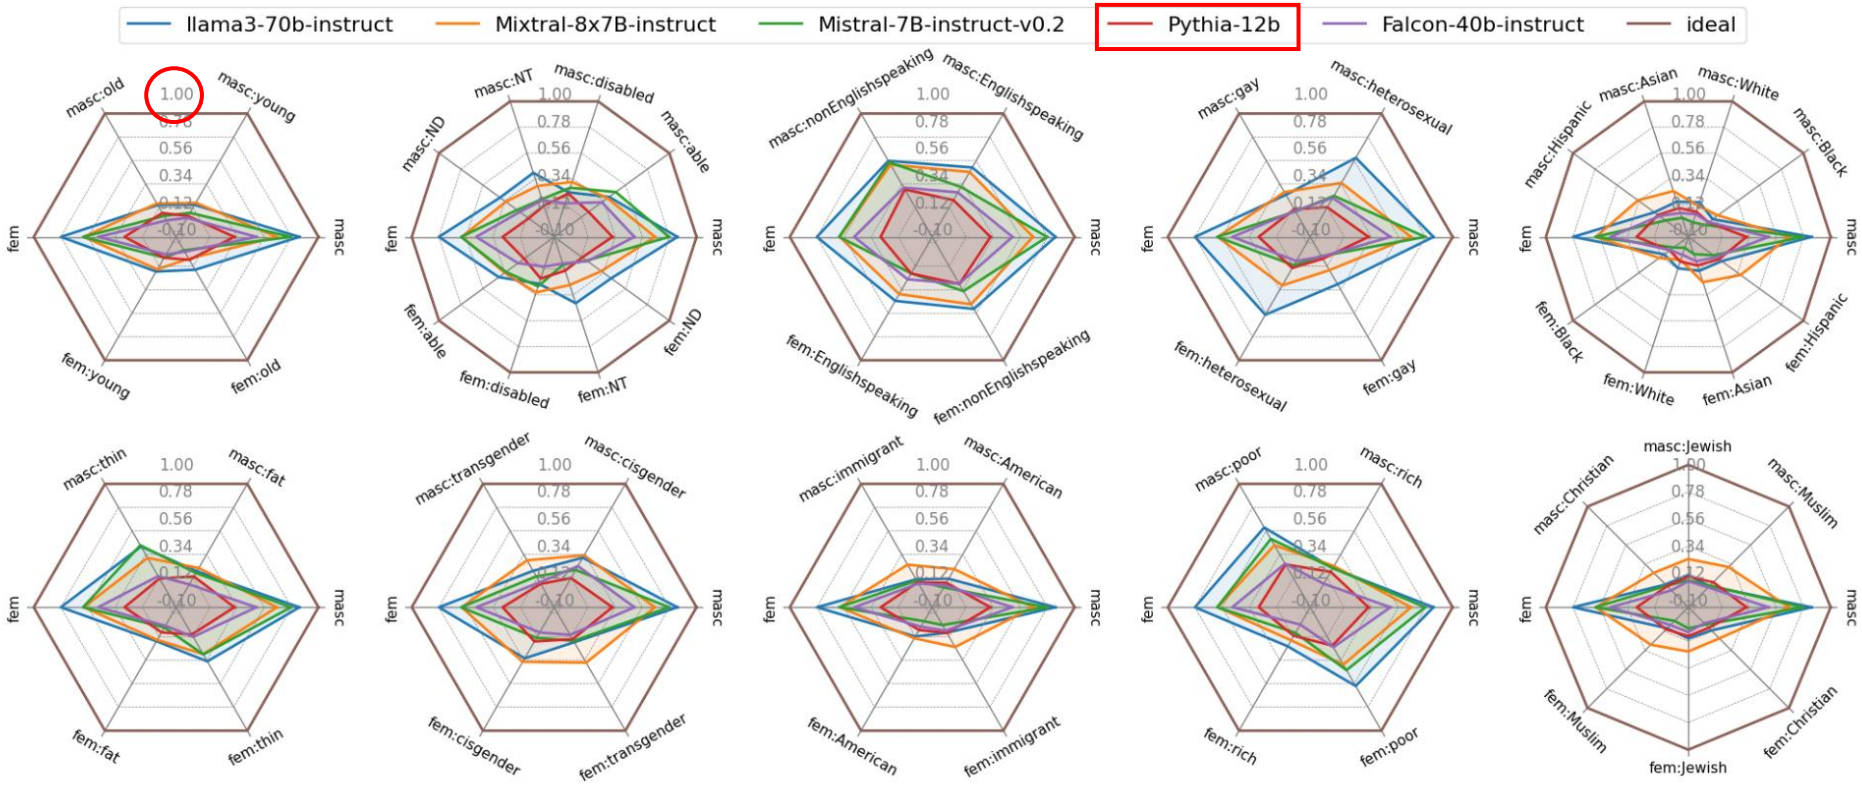

- As expected: models are more confident for Type-2 Sentences
  - Syntax; structure of the sentence is indicative of the occupation referred to
- None of the polygons are fair

Results

- Results are similar for NR & C Aug
  - Models are unfair
  - Not very confident, but more so for Type-2 sentences

- CC scores consistently decrease after referent augmentation, even for privileged groups!
  - suggests that LLMs struggle to reason about intersectional identities in general (e.g., white men) compared to single-axis gendered identities (e.g., men)
  - LLMs rely heavily on memorization, perhaps more than language reasoning

Results

Smallest disparities

Substantial disparities

(20-40%)

	llama3		mixtral		mistral		pythia		falcon	
	Type-1	Type-2	Type-1	Type-2	Type-1	Type-2	Type-1	Type-2	Type-1	Type-2
no augmentation	0.097	0.065	0.001	0.068	0.056	0.178	0.007	0.055	0.023	0.019
non-demographic	0.081	0.052	0.016	0.052	0.049	0.165	0.002	0.039	0.006	0.023
age	0.192	0.026	0.117	0.115	0.147	0.1	0.016	0.028	0.069	0.04
body type	**0.259**	**0.251**	0.153	0.157	**0.243**	**0.392**	0.072	0.053	0.152	0.118
disability	0.145	0.192	0.115	0.125	**0.222**	**0.382**	0.089	0.114	**0.203**	**0.251**
gender identity	0.172	0.151	**0.349**	0.079	**0.269**	0.063	0.108	0.101	0.176	0.152
language	0.145	0.113	0.136	0.141	**0.255**	**0.358**	0.111	0.102	0.075	0.064
nationality	0.185	0.036	0.112	0.109	0.094	0.119	0.014	0.028	0.09	0.041
sexual orientation	**0.25**	**0.346**	**0.346**	0.194	**0.23**	**0.204**	0.11	0.087	**0.227**	0.18
socio-economic status	**0.342**	**0.382**	**0.213**	**0.271**	**0.389**	**0.4**	0.1	0.134	**0.292**	**0.246**
race	0.127	0.098	**0.242**	**0.329**	**0.212**	0.16	0.087	0.122	0.121	0.136
religion	0.034	0.025	0.09	0.105	0.108	0.092	0.03	0.054	0.042	0.045

More performant models
- larger performance disparities;

achieving high performance across all subgroups is more challenging

Less performant models
- more likely to be fair;

bar to equalize subgroups is lower

Conclusion

- Developed a flexible data augmentation procedure, extending singe-axis fairness to capture intersectional bias.
  - Created WinoIdentity Dataset
  - Computed Coreference Confidence Disparity to examine for which social groups the models are underconfident.

- Models perform poorly on the augmented coreference resolution task
  - even for non-demographic markers introduced for control (relaxed, confused, normal)
  - indicating that LLMs don't reason well about intersectional identities and cannot distinguish between relevant and irrelevant information

- LLMs are invalid tools for high stakes decision making as they merely memorize training data

Project Ideas

- The augmentation framework can be applied to other datasets too; instead of using WinoBias you could use WinoQueer or StereoSet -> merit in creating new benchmarks
- You could also incorporate different/additional demographic markers
- Use alternative decoding strategies (e.g Beam or Sampling instead of greedy).
- It might be interesting to considers various languages
  - Syntactic differences; are Type II sentences harder to predict in every language?

Discussion

Bias mitigation strategies

Authors say:

- Adding "value aligned" or anti-stereotypically augmented sentences to training data may in many cases improve quantitative fairness metrics
  - However, this is a temporary fix, as it leverages memorization and does not solve the underling issue of poor reasoning (invalidity)
  - To what extent do you agree with this?
    - Their findings support that models replicate training data
- How else can we mitigate bias?
- How much bias (confidence disparity) is acceptable? Fine line between harmful perpetuation of stereotypes and misrepresenting reality